

0 Instructions (Total 40 pts)

Homework is due Tuesday, April 30, 2024 at 23:59pm Central Time. Please refer to <https://courses.grainger.illinois.edu/cs446/sp2024/homework/hw/index.html> for course policy on homeworks and submission instructions.

Version history.

- 1.0. Initial version.


1 Bellman Equation (7 pts)

1. (1 pts) What is the definition of the Q value function, $Q^\pi(s, a)$, for an infinite horizon with discount γ , reward function $R(s, a)$, stochastic transition dynamics $p(s' | s, a)$, with policy $\pi(a | s)$?
2. (2 pts) Derive the Bellman equation from this definition.
3. (1 pts) Given Q and transition (s, a, r, s') , what is the Q-learning update step (for discount γ and learning rate α)?
4. (1 pts) If the maximum reward is $R_{max} = \max_{s,a} R(s, a)$, with discount factor γ , what is the maximum value of the optimal Q, $\max_{s,a} Q^*(s, a)$?
5. (2 pts) Consider a system with two states S_1 and S_2 and two actions a_1 and a_2 . You perform actions and observe the rewards and transitions listed below. Each step lists the current state, reward, action and resulting transition as: $S_i; R = r; a_k : S_i \rightarrow S_j$. Perform Q-learning using a learning rate of $\alpha = 0.5$ and a discount factor of $\gamma = 0.5$ for each step by applying the formula from part (b). The Q-table entries are initialized to zero. Fill in the four tables below corresponding to the following four transitions. What is the final policy after having observed the four transitions?
 - (a) $S_1; R = -10; a_1 : S_1 \rightarrow S_1$
 - (b) $S_1; R = -10; a_2 : S_1 \rightarrow S_2$
 - (c) $S_2; R = 18.5; a_1 : S_2 \rightarrow S_1$
 - (d) $S_1; R = -10; a_2 : S_1 \rightarrow S_2$

2 Combination Lock (7 pts)

Consider an MDP with n states s_1, s_2, \dots, s_n , where there are 2 actions a_1, a_2 at each state. At each state s_i , $i < n$ action a_1 takes the agent to the s_{i+1} and for state s_n , action a_1 is a self-loop which keeps the agent in s_n . At any state, action a_2 takes the agent back to s_1 , except for the last state, which again has a self-loop. See the figure for an overall picture of the setting. $R(s_i, a_j)$ (reward of action a_j at state s_i) is 0, except for (s_n, a_1) , which has a value of 1. The agent takes one action at each step.

With *uniform random policy*, the agent at each state chooses one of the available actions uniformly at random. Now considering this *combination lock* MDP, answer the questions below. You need to show your work to receive full credit for each of the sections.



1. (2 pts) Compute the expected number of steps for the uniform random policy to go from state s_1 to state s_n .
2. (3 pts) Compute the formula for $Q(s_i, a_j)$, $\forall i, j$ for the uniform random policy considering a discounted reward setting with a discount factor of γ .
3. (2 pts) Now consider a new *greedy policy* using the $Q(s, a)$ function you computed in the previous part. Specifically, $\pi(s_i) = \operatorname{argmax}_a Q(s_i, a)$ (i.e., the agent, at each state, chooses the action with the highest value of the Q function computed in part (b)). Compute the new expected number of steps to get from state s_1 to s_k . Hint: how does $Q(s_i, a_1)$ compare to $Q(s_i, a_2)$?

3 Q-value initialization (7 pts)

In this question we will explore the effect of initialization on Q-learning. Let $Q_1(s, a) = 0$, $Q_2(s, a) = R_{max}$, where $R_{max} = \max_{s,a} R(s, a)$. Consider the greedy policy $\pi_i(s) = \operatorname{argmax}_a Q_i(s, a)$, where ties are broken at random (by selecting a random action among equal values).

Consider the MDP in the previous problem:

1. (2 pts) What is the expected number of steps for each policy to reach s_n starting from s_1 ?
2. (1 pts) Now assume that each time we explore a state transition (s, a, s', r) we update Q_i with a Bellman update, i.e.,

$$Q_i(s, a) = (1 - \alpha)Q_i(s, a) + \alpha(r + \gamma \max_{a'} Q_i(s', a'))$$

for $0 < \gamma < 1, 0 < \alpha < 1$.

Now what is the expected number of steps for π_1 to reach s_n starting from s_1 ?

3. (1 pts) Assuming we reset the agent (π_1) to s_1 once it reaches s_n , what is the expected number of steps before it reaches s_n again?
4. (1 pts) How can using a replay buffer help reduce the number of episodes before π_1 is optimal?
5. (2 pts) What are the minimum and the maximum number of steps for π_2 to reach s_n starting from s_1 (again assuming we update Q_2 each step)?

4 Policy Gradient (8 pts)

In policy gradient methods our goal is to evaluate the policy π directly rather than the value function. Let $J(\pi)$ be the expected value of the policy for some initial state distribution d_0 . In practice, π is a parameterized function, say π_θ .

- (3 pts) For trajectory $\tau = s_0, a_0, r_0, \dots, s_T, a_T, r_T$ the discounted sum of rewards is $R(\tau) = \sum_{i=0}^T \gamma^i r_i$ and $P^\pi(\tau)$ is the probability of τ under the policy, i.e., $P^\pi(\tau) = d_0(s_0) \prod_{i=1}^{T-1} \pi(a_i | s_i) P(s_{i+1} | s_i, a_i) \pi(a_T | s_T)$. Show that

$$\nabla_\theta J(\pi_\theta) = \mathbb{E}_{\tau \sim \pi_\theta} \left[R(\tau) \sum_{i=0}^T \nabla_\theta \log \pi_\theta(a_i | s_i) \right]$$

- (3 pts) An algorithm using the above gradient requires rolling out trajectories for each update. Let $d_i^\pi(s)$ be the probability the policy visits state s after i steps from the starting state distribution, and let $d^\pi(s)$ be the (discounted) probability that the policy visits s on some (infinite) trajectory, i.e., $d^\pi(s) = \sum_{i=0}^{\infty} \gamma^i d_i^\pi(s)$. Note this is called the stationary distribution, or discounted state occupancy. Show that

$$\nabla_\theta J(\pi_\theta) = \mathbb{E}_{s \sim d^\pi, a \sim \pi_\theta(s)} [Q^\pi(s, a) \nabla_\theta \log \pi_\theta(a | s)]$$

Hint: Notice that $J(\pi) = \mathbb{E}_{s \sim d_0} [V^\pi(s)]$, where V is the value function, and so $\nabla_\theta J(\pi_\theta) = \nabla_\theta [\mathbb{E}_{s \sim d_0} [V^{\pi_\theta}(s)]]$. Now how do you write $V^{\pi_\theta}(s)$ in terms of Q^{π_θ} ?

- (2 pts) Finally, show that

$$\nabla_\theta J(\pi_\theta) = \mathbb{E}_{s \sim d^\pi, a \sim \pi_\theta(s)} [(Q^\pi(s, a) - f(s)) \nabla_\theta \log \pi_\theta(a | s)]$$

for any function f . Many types of policy gradient depend on appropriate choice of f to reduce the variance of the estimated gradient.

5 Coding: Tabular Q-Learning (11 pts)

In this problem you will be implementing tabular Q learning for the Taxi-v3 environment.

1. (2 pts) Your first task is to investigate the OpenAI gym API and the Taxi-v3 environment. You can create the environment with

```
env = gym.make("Taxi-v3", render_mode="ansi")
```

- (a) What do states correspond to? How many possible states are there? What is the action space?
 - (b) What do `env.step()` and `env.reset()` return? How are states represented?
 - (c) What are the valid `render_modes` in `gym.make`? For Taxi-v3 what does each render mode do/return?
 - (d) Go digging in the open source documentation - how does `render()` convert the state representation returned by `step` and `reset` (hint: `self.s`) to the human readable features you described in part (a). Your answer should be a function call, `"self.function_name(self.s)"`.
2. (2 pts) Implement the `QAgent` class, more instructions are in the code. Report the `train_rewards.png` plot created by `q_learning()` for an experiment with $\alpha = 0.1, \gamma = 0.9, \epsilon = 0.1$ and 10,000 epochs.
 3. (2 pts) What happens if you initialize the Q table values to -10 (the minimum reward in the environment)? Why does it perform worse? Report the `train_moving_avg.png` plot (with -10 initial Q values) created by `q_learning()` for an experiment with $\alpha = 0.1, \gamma = 0.9, \epsilon = 0.1$ and 10,000 epochs.
 4. (1 pts) What percentage of the time does a uniform random policy ($\epsilon = 1.0$) get positive reward (success) in the environment across 10,000 epochs (printed for you as "Episode success rate")? Report the `train_rewards.png` plot for this uniform random policy.
 5. (1 pts) With initialization of -10 , report the `train_rewards.png` plot created by `q_learning()` for an experiment with $\alpha = 0.1, \gamma = 0.9$, 10000 epochs and ϵ of your choice which yields a model that successfully passes the evaluation (meaning the taxi delivers the person to the goal in the call to `"eval_agent()"`).
 6. (2 pts) Some actions lead to "no-op", for example moving right when there is a wall to the right. In the documentation for the Taxi-v3 environment they describe a method to only sample actions which lead to a change in state. Implement this; now, what percentage of the time does a uniform random policy ($\epsilon = 1.0$) get positive reward (success) in the environment across 10,000 epochs? Report the `train_rewards.png` plot.
 7. (1 pts) Again, with initialization of -10 AND never taking an action that leads to "no-op", report the `train_rewards.png` plot created by `q_learning()` for an experiment with $\alpha = 0.1, \gamma = 0.9$, 10000 epochs and ϵ the same as in part 5 which yields a model that successfully passes the evaluation.